



# Announcements

- HW1 due on Wednesday

# Recap: Keys and Foreign Keys

## Foreign Key

A **Key** is one or more attributes that **uniquely** identify a row.

A **Foreign Key** is one or more attrs that uniquely identify a row in *another table*.

# Recap: Keys and Foreign Keys

## Foreign Key

A **Key** is one or more attributes that **uniquely** identify a row.

A **Foreign Key** is one or more attrs that uniquely identify a row in *another table*.

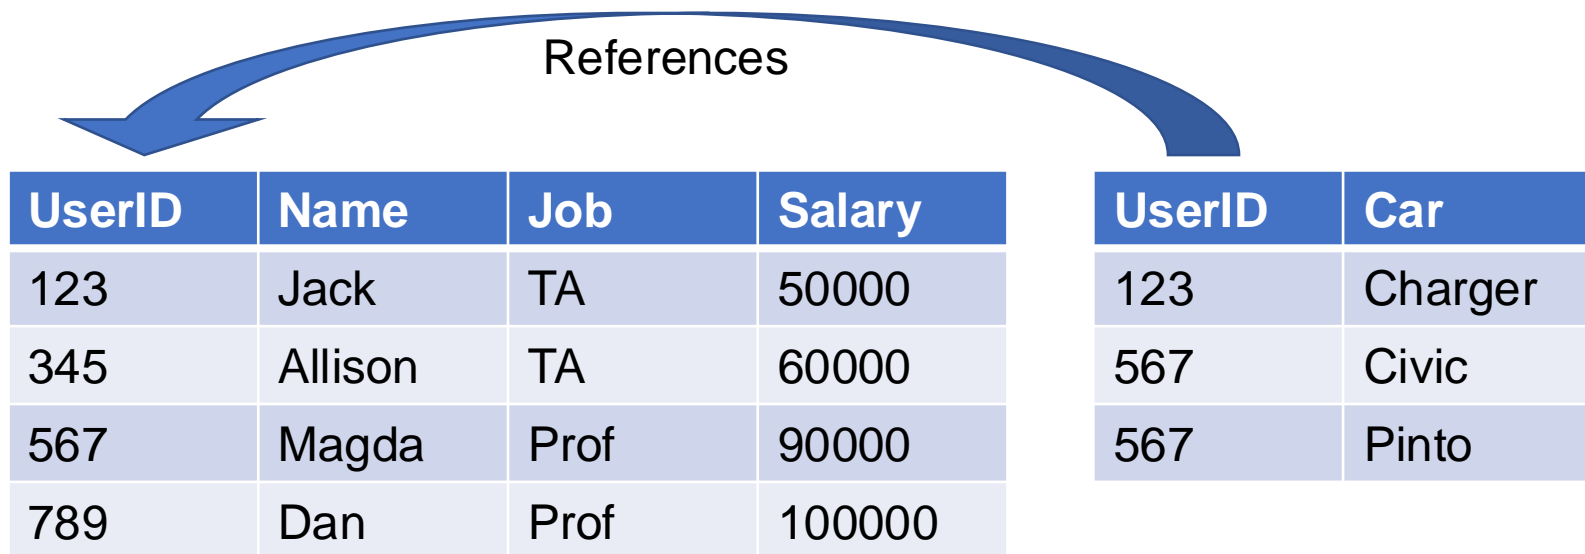
UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Recap: Keys and Foreign Keys

## Foreign Key

A **Key** is one or more attributes that **uniquely** identify a row.

A **Foreign Key** is one or more attrs that uniquely identify a row in *another table*.

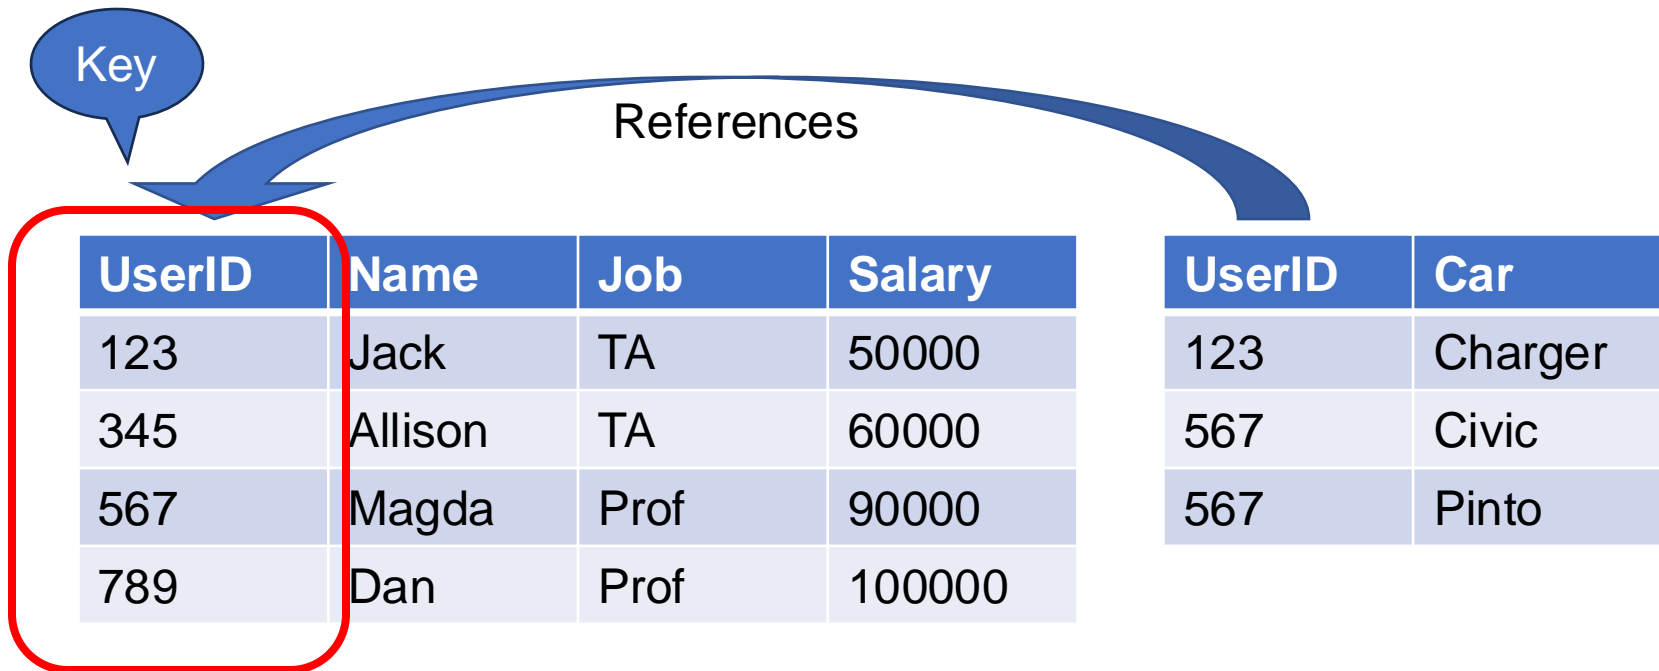


# Recap: Keys and Foreign Keys

## Foreign Key

A **Key** is one or more attributes that **uniquely** identify a row.

A **Foreign Key** is one or more attrs that uniquely identify a row in *another table*.

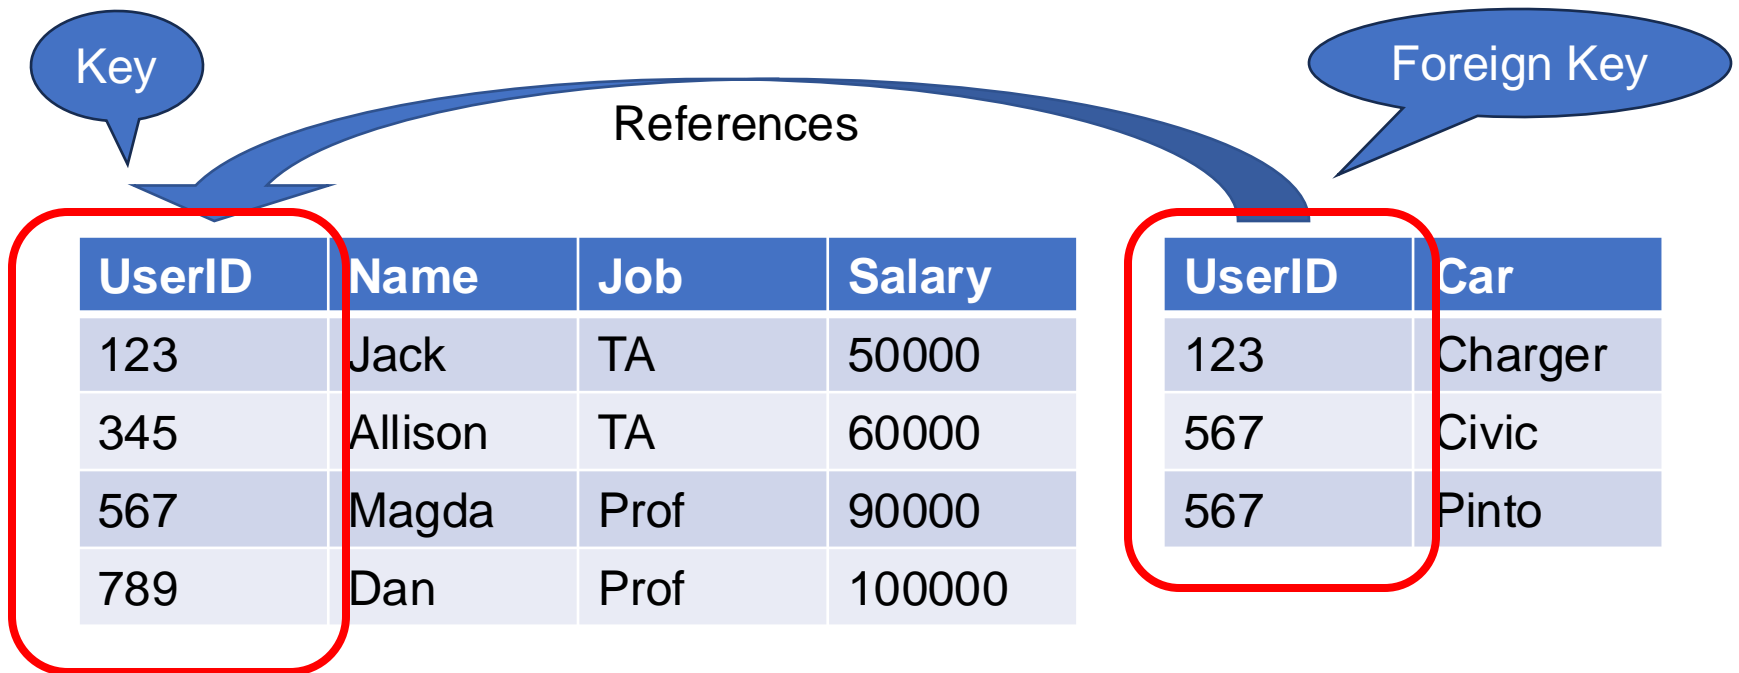


# Recap: Keys and Foreign Keys

## Foreign Key

A **Key** is one or more attributes that **uniquely** identify a row.

A **Foreign Key** is one or more attrs that uniquely identify a row in *another table*.



# Agenda

- Joins
- Nested Loop Semantics
- Self Joins
- Outer Joins



# Joins

- Joins link records from different tables.
- May use the key / foreign-key relationship, but may also use any other relationships

For each employee, find the cars that they drive

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```



Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join

For each TA, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
    and P.Job = 'TA';
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join

For each TA, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
       and P.Job = 'TA';
```



Name	Car
Jack	Charger

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto



# Join

For each TA, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
    and P.Job = 'TA';
```



Name	Car
Jack	Charger

and is a Boolean expression; let's review.

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Boolean Expression

- AND, OR, NOT
- Example: are these true or false?

$(5 < 7) \text{ or } (7 < 5)$

$(5 < 7) \text{ or } (5 < 8)$

$(5 < 7) \text{ and } (7 < 5)$

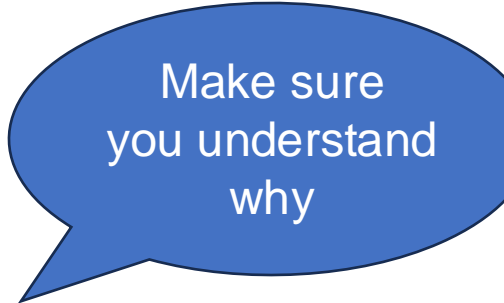
$(5 < 7) \text{ and not}(7 < 5)$

$(1 < 2) \text{ and } ((2 \neq 3) \text{ or } (4 < 3))$

# Boolean Expression

- AND, OR, NOT

- Example: are these true or false?



Make sure  
you understand  
why

$(5 < 7) \text{ or } (7 < 5)$

**TRUE**

$(5 < 7) \text{ or } (5 < 8)$

**TRUE**

$(5 < 7) \text{ and } (7 < 5)$

**FALSE**

$(5 < 7) \text{ and not}(7 < 5)$

**TRUE**

$(1 < 2) \text{ and } ((2 \neq 3) \text{ or } (4 < 3))$

**TRUE**

# Boolean Expression

In the WHERE clause: may use AND, OR, NOT

```
SELECT Name
FROM Payroll
WHERE Job = 'TA' or (Salary > 55000 and Salary < 95000);
```

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Boolean Expression

In the WHERE clause: may use AND, OR, NOT

```
SELECT Name
FROM Payroll
WHERE Job = 'TA' or (Salary > 55000 and Salary < 95000);
```

Name

Jack

Allison

Magda

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Boolean Expression

In the WHERE clause: may use AND, OR, NOT

```
SELECT Name
FROM Payroll
WHERE Job = 'TA' or (Salary > 55000 and Salary < 95000);
```

```
SELECT Name
FROM Payroll
WHERE Job = 'TA' and (Salary > 55000 and Salary < 95000);
```

Name

Jack

Allison

Magda

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Boolean Expression

In the WHERE clause: may use AND, OR, NOT

```
SELECT Name
FROM Payroll
WHERE Job = 'TA' or (Salary > 55000 and Salary < 95000);
```

```
SELECT Name
FROM Payroll
WHERE Job = 'TA' and (Salary > 55000 and Salary < 95000);
```

Name

Jack

Allison

Magda

Name

Allison

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

- When we use joins we often have multiple conditions in the `WHERE` clause: `and/or/not`
- Next: two ways to write the join



# Join: Two Ways to Write a Join

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join: Two Ways to Write a Join

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```

```
SELECT P.Name, R.Car  
FROM Payroll AS P JOIN Regist AS R  
ON P.UserID = R.UserID;
```

Means the  
same thing

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join: Two Ways to Write a Join

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
       and P.Job = 'TA';
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join: Two Ways to Write a Join

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
       and P.Job = 'TA';
```

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
       ON P.UserID = R.UserID
WHERE P.Job = 'TA';
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join: Two Ways to Write a Join

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
       and P.Job = 'TA';
```

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
       ON P.UserID = R.UserID
WHERE P.Job = 'TA';
```

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
       ON P.UserID = R.UserID
       and P.Job = 'TA';
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join: Two Ways to Write a Join

ON same as WHERE  
for now; but wait for it...

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
    and P.Job = 'TA';
```

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
    ON P.UserID = R.UserID
WHERE P.Job = 'TA';
```

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
    ON P.UserID = R.UserID
    and P.Job = 'TA';
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Discussion

- A join is often between a key and a foreign key
- But not always! Let's see some examples

# Join

```
-- find the cars they are driving
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto



# Join

```
-- find the cars they are driving
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

```
-- find the cars they are not driving
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID != R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

Payroll	UserID	Name	Job	Salary
	123	Jack	TA	50000
	345	Allison	TA	60000
	567	Magda	Prof	90000
	789	Dan	Prof	100000

Regist	UserID	Car
	123	Charger
	567	Civic
	567	Pinto

# Join

```
-- find the cars they are driving
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

```
-- find the cars they are not driving
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID != R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto



Name	Car
Jack	Civic
Jack	Pinto
Allison	Charger
Allison	...
...	

Payroll	UserID	Name	Job	Salary
	123	Jack	TA	50000
	345	Allison	TA	60000
	567	Magda	Prof	90000
	789	Dan	Prof	100000

Regist	UserID	Car
	123	Charger
	567	Civic
	567	Pinto

# Join

```
-- find the cars they are driving
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

```
-- find the cars they are not driving
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID != R.UserID;
```

```
-- find WHAT??
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

Name	Car
Jack	Civic
Jack	Pinto
Allison	Charger
Allison	...
...	

Payroll				
	UserID	Name	Job	Salary
	123	Jack	TA	50000
	345	Allison	TA	60000
	567	Magda	Prof	90000
	789	Dan	Prof	100000

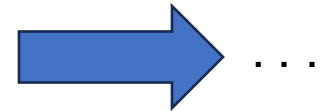
Regist		
	UserID	Car
	123	Charger
	567	Civic
	567	Pinto

# Join

```
-- find the cars they are driving
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

```
-- find the cars they are not driving
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID != R.UserID;
```

```
-- find WHAT??
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

Name	Car
Jack	Civic
Jack	Pinto
Allison	Charger
Allison	...
...	

Payroll	UserID	Name	Job	Salary
	123	Jack	TA	50000
	345	Allison	TA	60000
	567	Magda	Prof	90000
	789	Dan	Prof	100000

Regist	UserID	Car
	123	Charger
	567	Civic
	567	Pinto

# Discussion

- FROM clause: several table names
- WHERE clause: some condition on these tables
- **Q:** What does it mean?
- **A:** For-Each semantics (Nested Loop Semantics)!

# Nested Loop Semantics (again!)

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto



# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```

How do we  
algorithmically get  
our results?

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

# Nested-Loop Semantics


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```


```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000


UserID	Car
123	Charger
567	Civic
567	Pinto




Name	Car
------	-----

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000




UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger


```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
Jack	Charger


```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```



# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```



# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Final answer

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```



Key / Foreign-key join

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```



Key / Foreign-key join

```
for each row1 in Payroll:  
  for each row2 in Regist:  
    if (row1.UserID = row2.UserID):  
      output (row1.Name, row2.Car)
```

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R;
```



Cross product

```
for each row1 in Payroll:  
  for each row2 in Regist:  
    output (row1.Name, row2.Car)
```

# Summary: Nested-Loop Semantics

- FROM clause contains tables  $T_1, T_2, T_3, \dots$
- WHERE clause contains `condition`
- SELECT clause contains `attr1, attr2, \dots`

```
for each r1 in T1:
  for each t2 in T2:
    for each t3 in T3:
      ...
      if (condition):
        output (attr1, attr2, ...)
```

# Self-Joins

# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto



# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

Intended  
answer

Name	Car1	Car2
Magda	Civic	Pinto

# Self Joins

Find all people who drive a ~~Civic~~ and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto



Let's start  
with Pinto...

# Self Joins

Find all people who drive a ~~Civic~~ and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
        R.Car = 'Pinto';
```



Let's start  
with Pinto...

# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      R.Car = 'Civic' AND
      R.Car = 'Pinto';
```



Now both

# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      R.Car = 'Civic' AND
      R.Car = 'Pinto';
```

Now both

Will this work?

# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
  FROM Payroll AS P, Regist AS R
 WHERE P.UserID = R.UserID AND
        R.Car = 'Civic' AND
        R.Car = 'Pinto';
```

Will this work?  
Nope, returns  
the empty set.

# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      (R.Car = 'Civic' OR
       R.Car = 'Pinto');
```

Is this better?

# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
  FROM Payroll AS P, Regist AS R
 WHERE P.UserID = R.UserID AND
       (R.Car = 'Civic' OR
        R.Car = 'Pinto');
```

Is this better?

Nope, it returns  
both Jack and Magda.



# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      (R.Car = 'Civic' OR
       R.Car = 'Pinto');
```

Discuss with the people around you how you would solve this.

# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

```
SELECT P.Name, R1.Car as Car1, R2.Car as Car2
FROM Payroll AS P, Regist AS R1, Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```


# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

```
SELECT P.Name, R1.Car as Car1, R2.Car as Car2
FROM Payroll AS P, Regist AS R1, Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```



Name	Car1	Car2
Magda	Civic	Pinto

# Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
123	Pinto
123	Tesla
567	Civic
567	Pinto

The person we look for  
must drive TWO cars

```
SELECT P.Name, R1.Car as Car1, R2.Car as Car2
FROM Payroll AS P, Regist AS R1, Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```

Name	Car1	Car2
Magda	Civic	Pinto

# Self Joins

- When a relation occurs twice in the FROM clause we call it a “self-join”
- If we have a self-join, we must use table aliases; Otherwise, the attribute names are ambiguous

# Outer Joins

# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P
      JOIN Regist AS R
      ON P.UserID = R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P
      JOIN Regist AS R
      ON P.UserID = R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

Allison, Dan  
are missing

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto



# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P
      LEFT OUTER JOIN Regist AS R
      ON P.UserID = R.UserID;
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P
      LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto
Allison	NULL
Dan	NULL

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P
      LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto
Allison	NULL
Dan	NULL

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

NULL means  
“unknown” or  
“missing”

# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P
      LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto
Allison	NULL
Dan	NULL

Left outer join:

1. Perform the join with the ON clause
2. Add all missing tuples from LEFT
3. Check the WHERE clause (if present)

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Join

For each employee, find the cars that they drive

```
SELECT P.Name, R.Car
FROM Payroll AS P
      LEFT OUTER JOIN Regist AS R
      ON P.UserID = R.UserID;
```



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto
Allison	NULL
Dan	NULL

Left outer join:

1. Perform the join with the ON clause
2. Add all missing tuples from LEFT
3. Check the WHERE clause (if present)

ON, WHERE differ  
(next lecture)

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

# Outer Joins

- LEFT OUTER JOIN
  - Add missing tuples from the LEFT
- RIGHT OUTER JOIN
  - Add missing tuples from the RIGHT
- FULL OUTER JOIN
  - Add missing tuples from both

# Outer Joins

- LEFT OUTER JOIN
  - Add missing tuples from the LEFT

Useful, especially  
for aggregates  
(next lecture)

- RIGHT OUTER JOIN
  - Add missing tuples from the RIGHT

- FULL OUTER JOIN
  - Add missing tuples from both

Rarely  
used