

# Data Management for Data Science

## SQL Basics

Paul G. Allen School of Computer Science and Engineering  
University of Washington, Seattle

# Announcements

- HW 1 released – due Friday 1/12 at **11pm**
  - Submitted via gradescope
  - Try to do HW 1 setup today (should take ~5-10 minutes)
  - Yesterday's section demo should be really useful!
  - The demo and all other section materials are on the course website

# Recap - Data Models

The 3 parts of any data model

- **Instance**
  - The actual **data**
- **Schema**
  - A **description** of what data is being stored
- **Query Language**
  - How to retrieve and manipulate data

**Medical Records**

PatientID	Name	Status	Notes
123	Alex	Healthy?	...
345	Bob	Critical	...

# Recap - The Relational Model

- Flat tables, static and typed attributes, etc.
  - “It’s a spreadsheet with rules”

**Table/  
Relation**

**Columns/Attributes/Fields**

**Rows/  
Tuples/  
Records**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Recap - The Relational Model

But how is this data ACTUALLY stored?

**Payroll**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Recap - The Relational Model

But how is this data ACTUALLY stored?

**Payroll**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Don't know. Don't care.

**Physical Data Independence**

# Structured Query Language - SQL

Alright, I have data and a schema.  
How do I access it?

# Structured Query Language - SQL

- **Declarative** query language
  - Tell the computer what you want, not how to get it
- Languages like Java/Python are procedural
- Declarative query language allows **physical data independence**



# Basic SQL query

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



```
SELECT *  
FROM Payroll;
```

# Basic SQL query

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000




```
SELECT *  
FROM Payroll;
```

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Basic SQL query

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000




```
SELECT P.Name, P.UserID
FROM Payroll AS P
WHERE P.Job = 'TA';
```

# Basic SQL query

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000




Name	UserID
Jack	123
Allison	345

```
SELECT P.Name, P.UserID
FROM Payroll AS P
WHERE P.Job = 'TA';
```

# Basic SQL query

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



Name	UserID
Jack	123
Allison	345

```
SELECT P.Name, P.UserID  
FROM Payroll AS P  
WHERE P.Job = 'TA';
```

“Payroll AS P” makes P an alias.  
This lets us specify that the  
attributes come from Payroll

Wait!

What actually  
happens when we  
execute the SQL  
query?

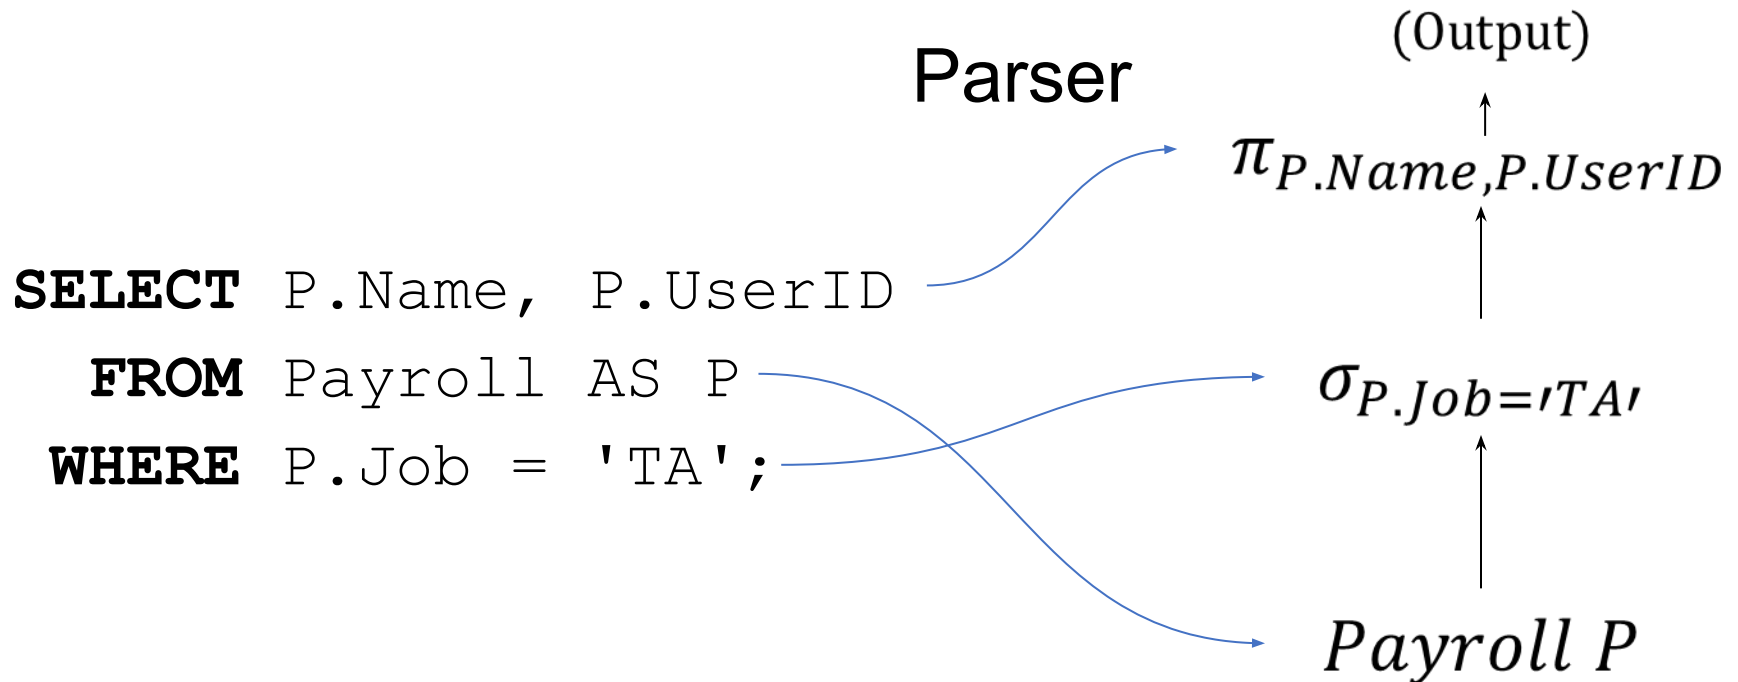
# Database Internals

- Code has to boil down to instructions at some point
- Relational Database Management Systems (RDBMSs) use **Relational Algebra** (RA)

```
SELECT P.Name, P.UserID  
      FROM Payroll AS P  
      WHERE P.Job = 'TA';
```

# Database Internals

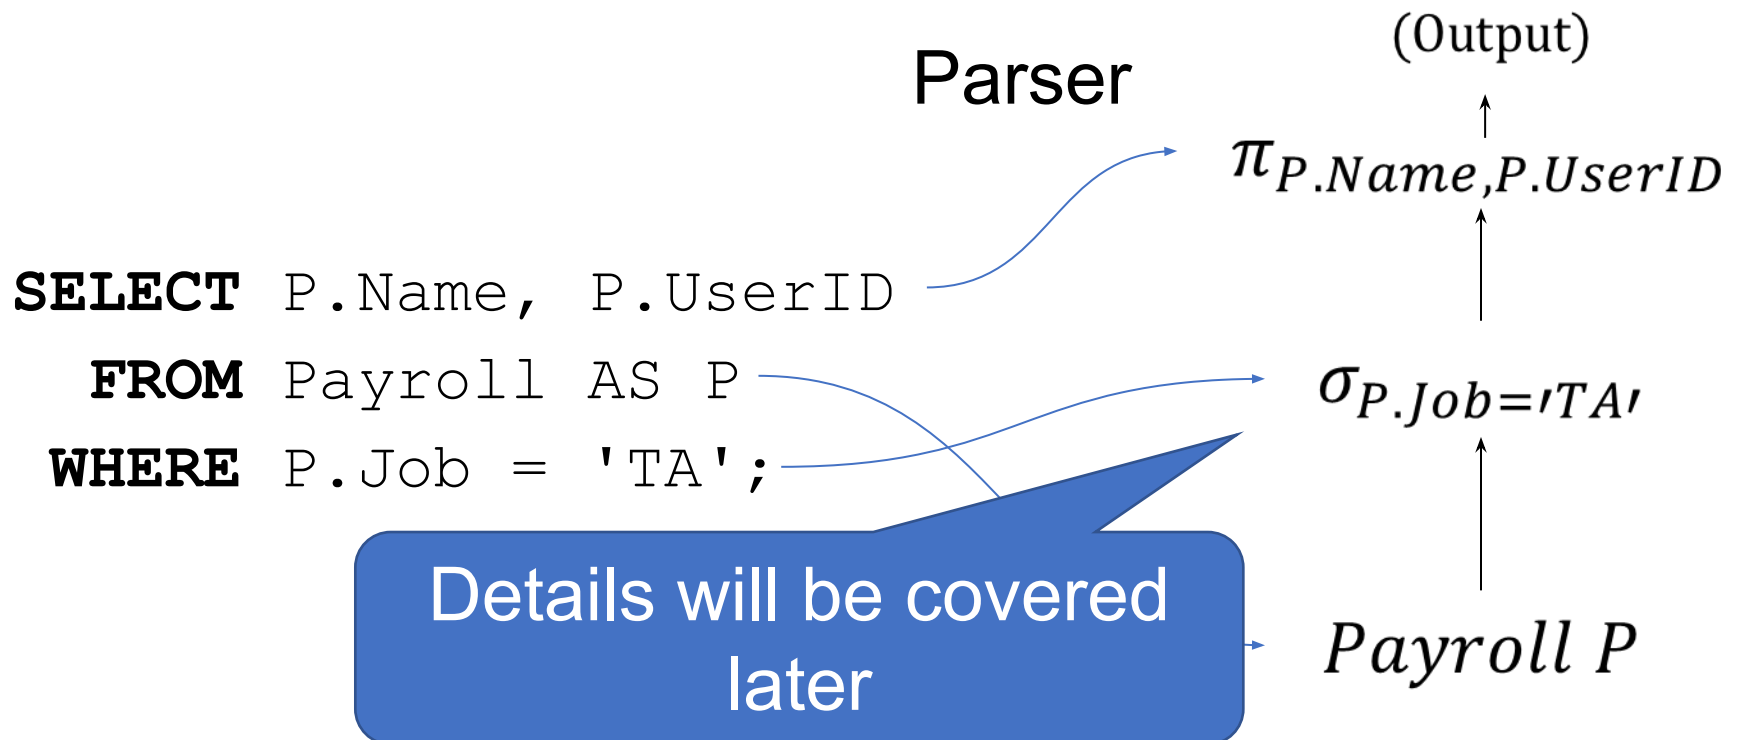
- Code has to boil down to instructions at some point
- Relational Database Management Systems (RDBMSs) use **Relational Algebra** (RA).





# Database Internals

- Code has to boil down to instructions at some point
- Relational Database Management Systems (RDBMSs) use **Relational Algebra** (RA).




# Database Internals

- It's important to define the semantics (meaning) of a query

```
SELECT P.Name, P.UserID  
FROM Payroll AS P  
WHERE P.Job = 'TA';
```

For-each semantics



# Database Internals

- It's important to define the semantics (meaning) of a query


```
SELECT P.Name, P.UserID  
FROM Payroll AS P  
WHERE P.Job = 'TA';
```

## For-each semantics

for each row in P:

if (row.Job == 'TA'):

output (row.Name, row.UserID)



# Database Internals

- It's important to define the semantics (meaning) of a query

```
SELECT P.Name, P.UserID  
FROM Payroll AS P  
WHERE P.Job = 'TA';
```

(Output)

$\pi_{P.Name, P.UserID}$

$\sigma_{P.Job='TA'}$

*Payroll P*

```
graph BT; P["Payroll P"] --> S["σP.Job='TA'"]; S --> P2["πP.Name, P.UserID"]; P2 --> O["(Output)"];
```

Tuples “flow” up the query plan, getting filtered and modified

# For-Each Semantics

## Payroll


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

for each row in P:

if (row.Job == 'TA'):

output (row.Name, row.UserID)

Job == 'TA'?



Name	UserID
------	--------

```
SELECT P.Name, P.UserID
FROM Payroll AS P
WHERE P.Job = 'TA';
```

# For-Each Semantics

## Payroll


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

for each row in P:

if (row.Job == 'TA'):

output (row.Name, row.UserID)

Job == 'TA'?



Name	UserID
Jack	123

```
SELECT P.Name, P.UserID
FROM Payroll AS P
WHERE P.Job = 'TA';
```

# For-Each Semantics

## Payroll


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

for each row in P:

if (row.Job == 'TA'):

output (row.Name, row.UserID)

Job == 'TA'?



Name	UserID
Jack	123

```
SELECT P.Name, P.UserID
FROM Payroll AS P
WHERE P.Job = 'TA';
```

# For-Each Semantics

## Payroll


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

for each row in P:

if (row.Job == 'TA'):

output (row.Name, row.UserID)

Job == 'TA'?



Name	UserID
Jack	123
Allison	345

```
SELECT P.Name, P.UserID
FROM Payroll AS P
WHERE P.Job = 'TA';
```



# For-Each Semantics

## Payroll


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

for each row in P:

if (row.Job == 'TA'):

output (row.Name, row.UserID)

Job == 'TA'?



Name	UserID
Jack	123
Allison	345

```
SELECT P.Name, P.UserID
FROM Payroll AS P
WHERE P.Job = 'TA';
```

# For-Each Semantics

## Payroll


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

for each row in P:

if (row.Job == 'TA'):

output (row.Name, row.UserID)

Job == 'TA'?



Name	UserID
Jack	123
Allison	345


```
SELECT P.Name, P.UserID
FROM Payroll AS P
WHERE P.Job = 'TA';
```

# For-Each Semantics

## Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

```
for each row in P:  
  if (row.Job == 'TA'):  
    output (row.Name, row.UserID)
```



Name	UserID
Jack	123
Allison	345

```
SELECT P.Name, P.UserID  
FROM Payroll AS P  
WHERE P.Job = 'TA';
```

# Recap – SQL and RA

## ■ SQL

(Next few lectures)

- “What data do I want”


## ■ RA

(After SQL)

- “How do I get the data”

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

```
SELECT P.Name, P.UserID
FROM Payroll AS P
WHERE P.Job = 'TA';
```



Name	UserID
Jack	123
Allison	345

# What's Next?

- Creating tables
- Keys □ Identification
- Foreign Keys □ Relationships
- Joins in SQL and RA
  - Inner joins
  - Outer joins
  - Self joins

# Create Table Statement

Payroll(UserId, Name, Job, Salary)



```
CREATE TABLE Payroll (  
    UserID INT,  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT);
```

# Data Types

- Each attribute has a type.
  - Examples types:
    - Strings: CHAR(20), VARCHAR(50), TEXT
    - Numbers: INT, SMALLINT, FLOAT
    - MONEY, DATETIME, ...
    - Few more that are DBMS specific
  - Statically and strictly enforced

# Data Types

- Generally you will use:
  - **VARCHAR(N)** for strings where **N** is the maximum character length
    - Generally set this to as large as you need, like 256 or 1000.
  - **INT**, **FLOAT** for numbers (**INTEGER** works in SQLite)
  - **DATETIME** for dates
    - Can use **VARCHAR(N)** in SQLite



# Create Table Statement

Payroll(**UserId**, Name, Job, Salary)



```
CREATE TABLE Payroll (  
    UserID INT,  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT);
```

# Create Table Statement

Payroll(**UserId**, Name, Job, Salary)



```
CREATE TABLE Payroll (  
    UserID INT,  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT);
```

Everything is case-insensitive, but having your own guidelines is useful for readability

# Keys

## Key

A **Key** is one or more attributes that **uniquely** identify a row.

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Keys

## Key

A **Key** is one or more attributes that **uniquely** identify a row.

Definitely not a key

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Keys

## Key

A **Key** is one or more attributes that **uniquely** identify a row.

Good candidate  
for a key

Definitely not a key

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Keys

## Key

A **Key** is one or more attributes that **uniquely** identify a row.

Is this a good candidate for a key?

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Keys

## Key

A **Key** is one or more attributes that **uniquely** identify a row.

Is this a good candidate for a key?

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000
913	Peter	TA	60000

# Keys

## Key

A **Key** is one or more attributes that **uniquely** identify a row.

Data comes from the real world  
so models ought to reflect that

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000
913	Peter	TA	60000



```
CREATE TABLE Payroll (  
    UserID INT,  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT);
```

Payroll(UserId, Name, Job, Salary)

# Keys

Unique Identifier

```
CREATE TABLE Payroll (  
  UserID INT,  
  Name VARCHAR(100),  
  Job VARCHAR(100),  
  Salary INT);
```

Payroll(UserId, Name, Job, Salary)

# Keys

Unique Identifier

```
CREATE TABLE Payroll (  
  UserID INT PRIMARY KEY,  
  Name VARCHAR(100),  
  Job VARCHAR(100),  
  Salary INT);
```

Payroll(UserId, Name, Job, Salary)

# Keys

```
CREATE TABLE Payroll (  
    UserID INT,  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT,  
    PRIMARY KEY (UserId);
```

Can also define the  
PK on a new line

Payroll(UserId, Name, Job, Salary)

# Keys of more than one attribute

Sometimes no single attribute is unique, but combinations of attributes are a unique key for the table.

```
CREATE TABLE Payroll (  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT,  
    PRIMARY KEY (Name, Job));
```

Must use the PK definition on a new line for multi-attribute keys

# Keys of more than one attribute

Sometimes no single attribute is unique, but combinations of attributes are a unique key for the table.

Must use the PK definition on a new line for multi-attribute keys

```
CREATE TABLE Payroll (  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT,  
    PRIMARY KEY (Name, Job));
```

Here the combination of Name and Job are unique  
e.g. only one "Eden, Professor"  
but some "Eden, TA" or "Ryan, Professor" can exist

Payroll(Name, Job, Salary)

# A little extra SQL

- **ORDER BY** – Orders result tuples by specified attributes (default ascending)

```
SELECT P.Name, P.UserID  
      FROM Payroll AS P  
      WHERE P.Job = 'TA'  
      ORDER BY P.Salary, P.Name;
```

- **DISTINCT** – Deduplicates result tuples

```
SELECT DISTINCT P.Job  
      FROM Payroll AS P  
      WHERE P.Salary > 70000;
```

# Foreign Keys

- Databases can hold multiple tables
- How do we capture relationships *between* tables?

**Payroll**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Regist**

UserID	Car
123	Charger
567	Civic
567	Pinto



# Foreign Keys

- Databases can hold multiple tables
- How do we capture relationships *between* tables?

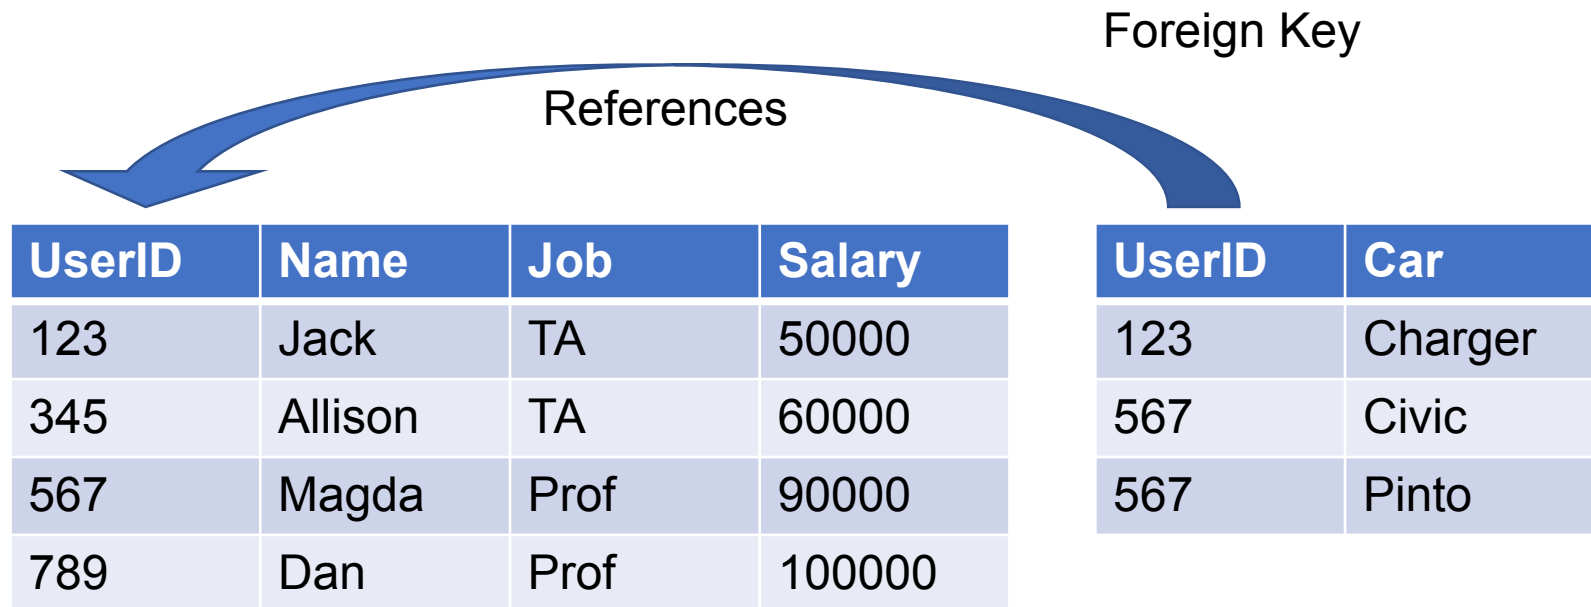
Foreign Key  
UserID

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

# Foreign Keys

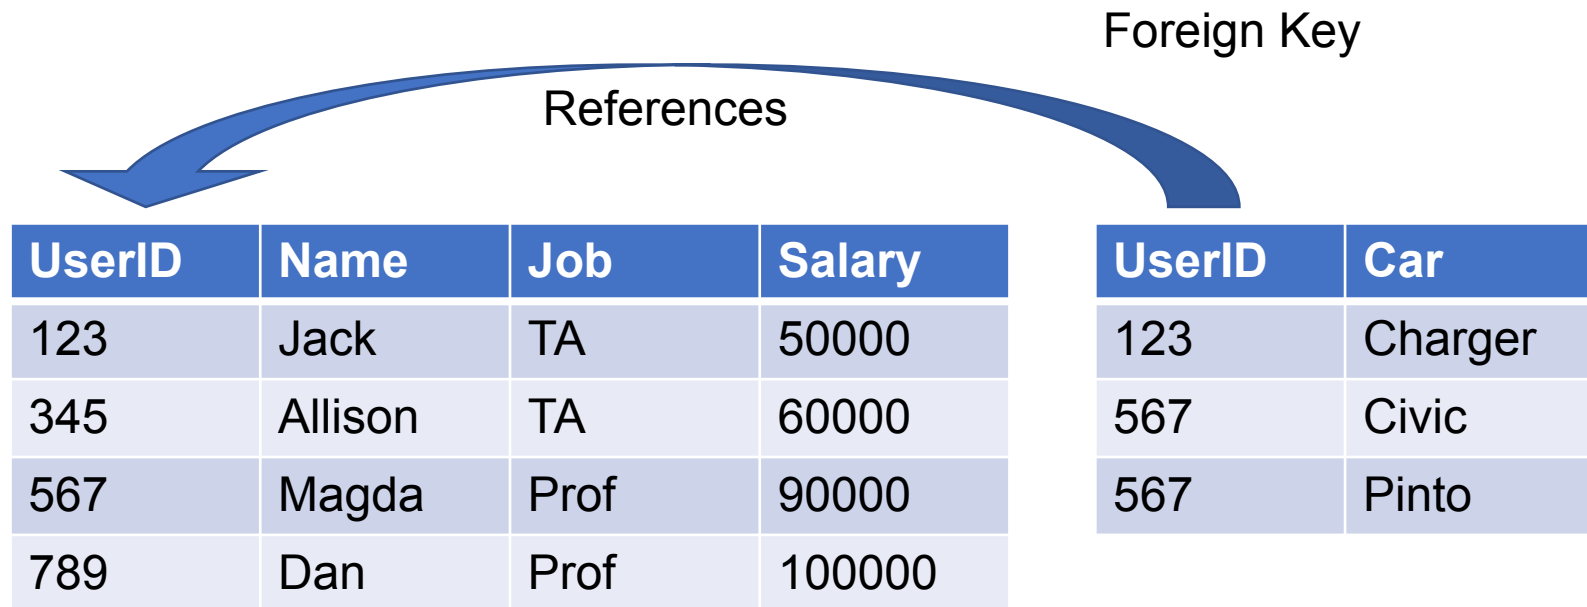
- Databases can hold multiple tables
- How do we capture relationships *between* tables?



# Foreign Keys

## Foreign Key

A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.



# Foreign Keys

## Foreign Key

A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.

Is this valid?

References

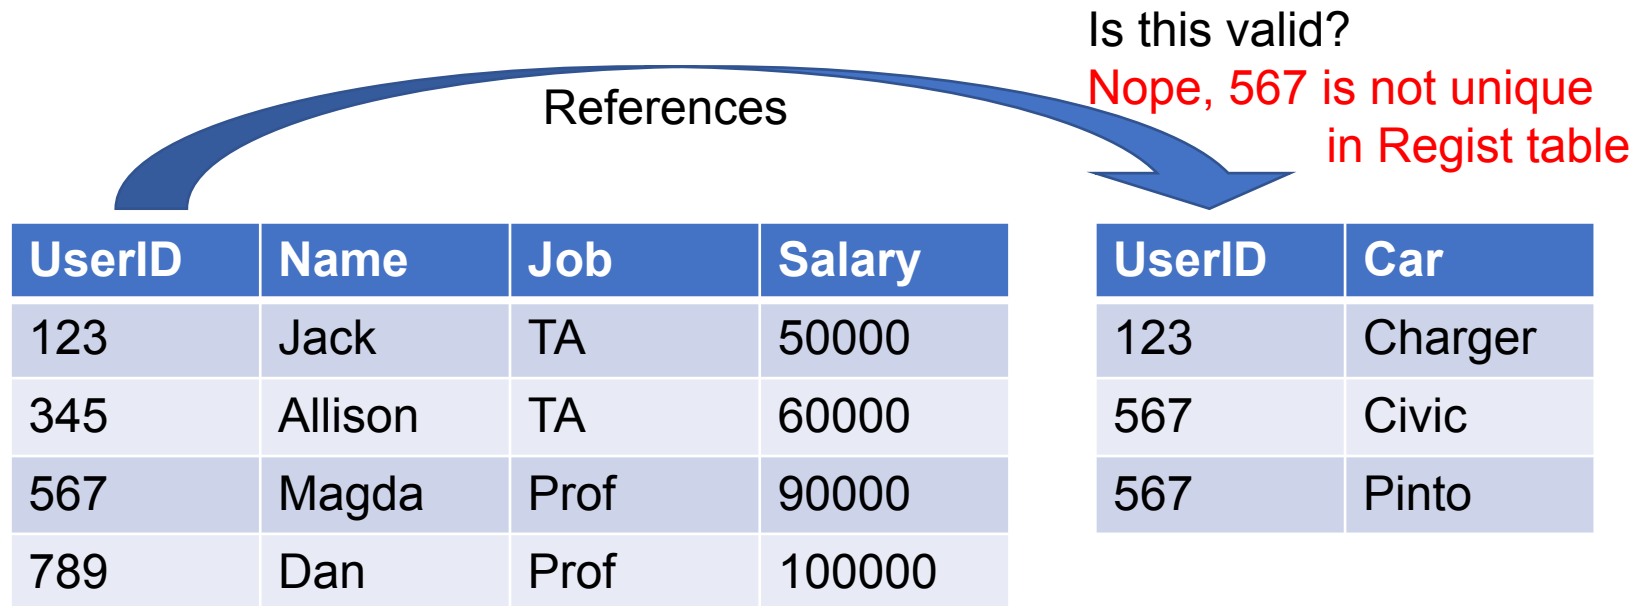
UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

# Foreign Keys

## Foreign Key

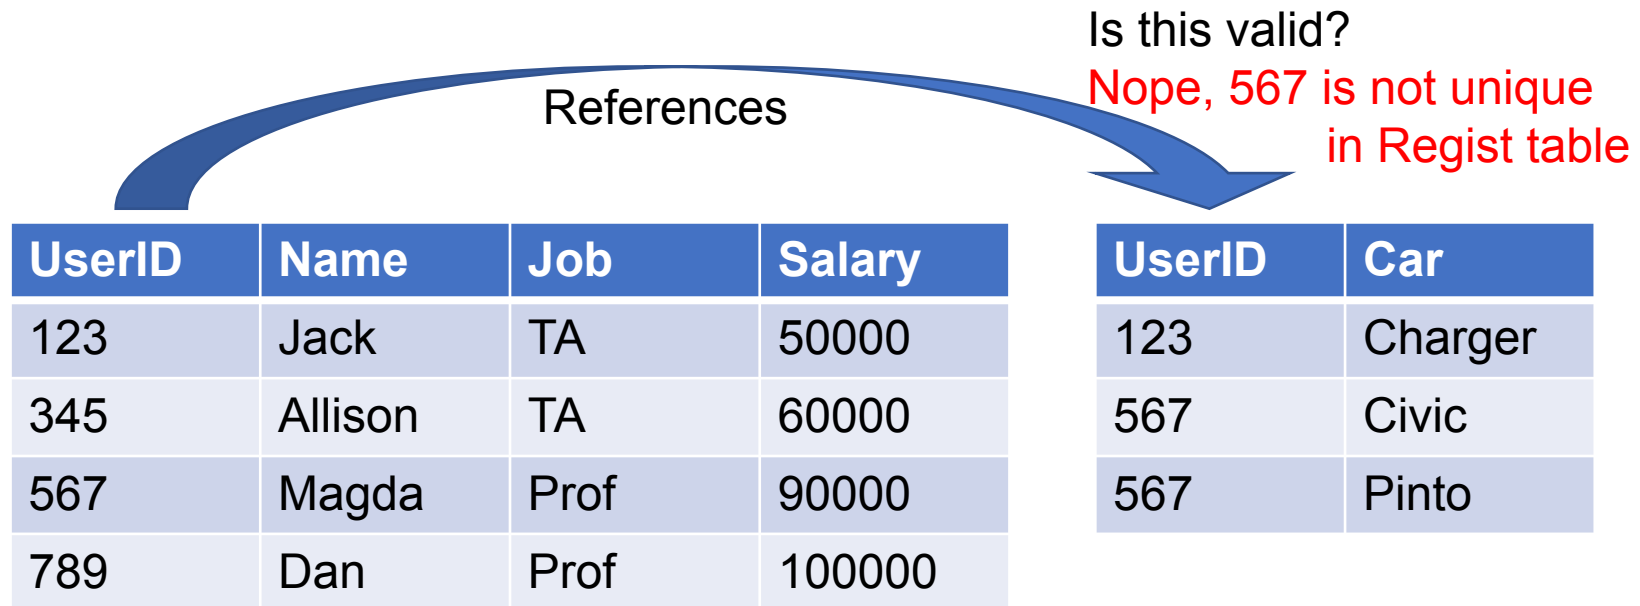
A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.



# Foreign Keys

## Foreign Key

A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.



Foreign keys must reference (point to) a unique attribute, almost always a primary key

# Foreign Keys

```
CREATE TABLE Payroll (  
    UserID INT PRIMARY KEY,  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT);
```

Payroll(UserId, Name, Job, Salary)

```
CREATE TABLE Regist (  
    UserID INT,  
    Car VARCHAR(100));
```

Regist(UserId, Car)

# Foreign Keys

```
CREATE TABLE Payroll (  
    UserID INT PRIMARY KEY,  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT);
```

Payroll(UserId, Name, Job, Salary)

```
CREATE TABLE Regist (  
    UserID INT REFERENCES Payroll,  
    Car VARCHAR(100));
```

Regist(UserId, Car)



# Foreign Keys

```
CREATE TABLE Payroll (  
  UserID INT PRIMARY KEY,  
  Name VARCHAR(100),  
  Job VARCHAR(100),  
  Salary INT);
```

```
CREATE TABLE Regist (  
  UserID INT REFERENCES Payroll(UserID),  
  Car VARCHAR(100));
```

or, when attribute name is the same:

```
CREATE TABLE Regist (  
  UserID INT REFERENCES Payroll,  
  Car VARCHAR(100));
```

Payroll(UserId, Name, Job, Salary)

Regist(UserId, Car)

# Foreign Keys

Alternatively, if your foreign key is also more than one attribute:

```
CREATE TABLE Payroll (  
    UserID INT,  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT,  
    PRIMARY KEY (UserID,  
        Name)  
);
```

Payroll(UserID, Name, Job, Salary)

```
CREATE TABLE Regist (  
    UserID INT,  
    Name VARCHAR(100),  
    Car VARCHAR(100),  
    FOREIGN KEY (UserID, Name)  
        REFERENCES Payroll);
```

Regist(UserID, Name, Car)

# The Relational Model Revisited

- More complete overview of the Relational Model:
  - Database □ collection of tables
  - All tables are flat
  - Keys uniquely ID rows
  - Foreign keys act as a “semantic pointer”
  - **Physical data independence**

# Joins

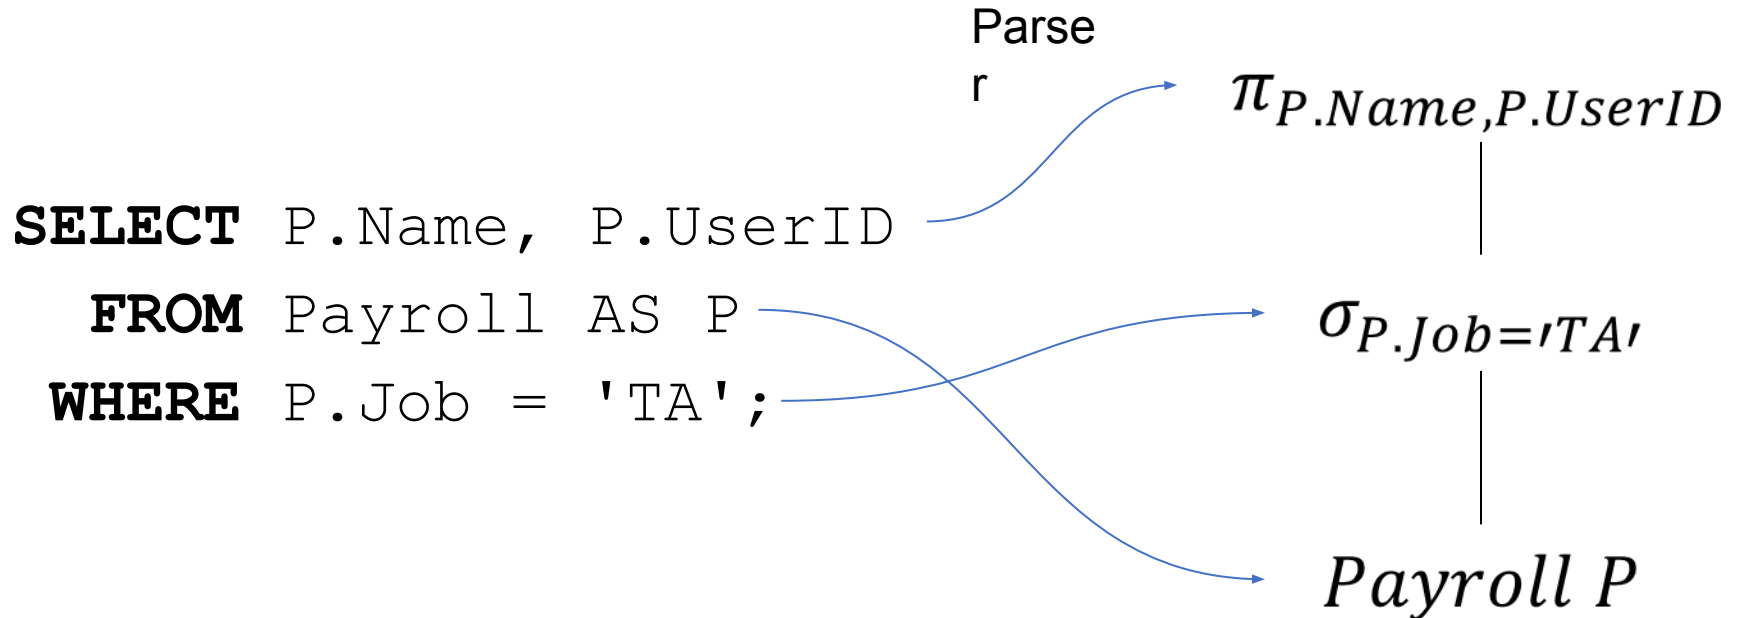
- Foreign keys are able to *describe* a relationship between tables
- Joins are able to realize combinations of data

# RA Equivalencies

So far we haven't discussed equivalent RA trees.  
But all joins can be parsed directly into a “join tree”

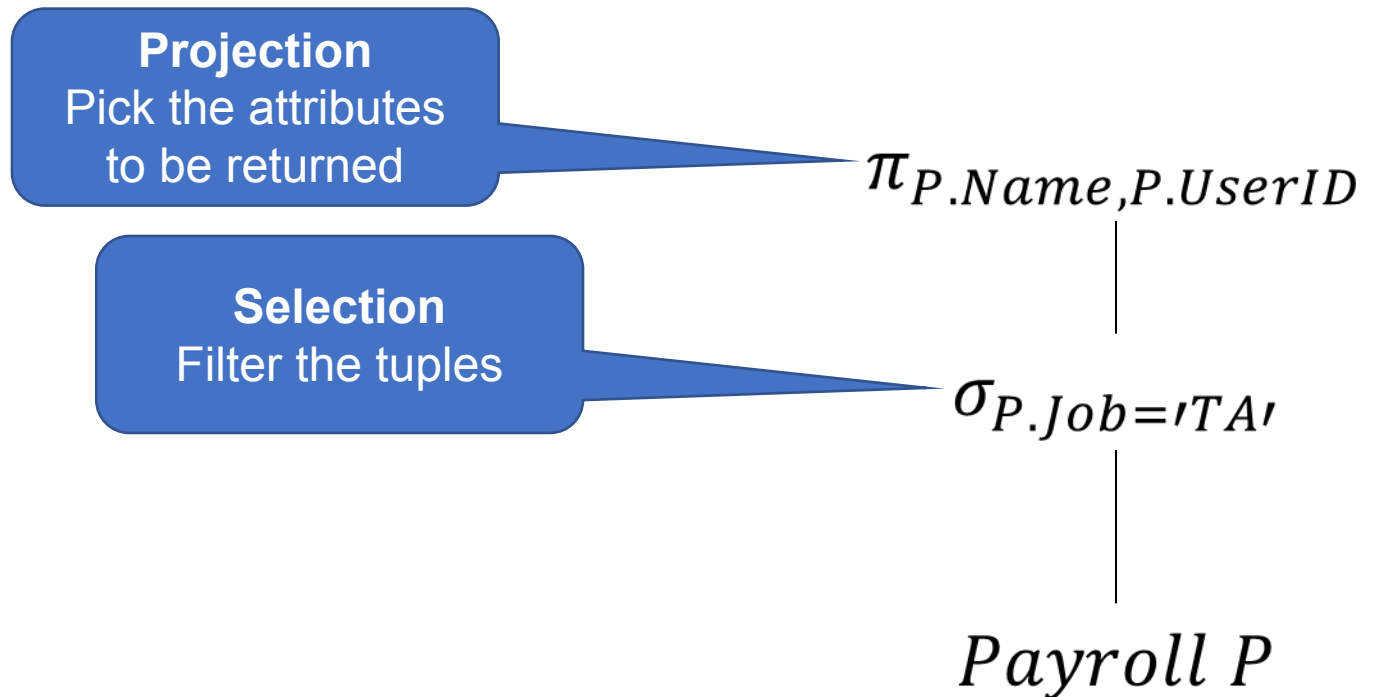
# RA Equivalencies

So far we haven't discussed equivalent RA trees.  
But all joins can be parsed directly into a “join tree”




# RA Equivalencies

So far we haven't discussed equivalent RA trees.  
But all joins can be parsed directly into a “join tree”



# RA Equivalencies

```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```






# RA Equivalencies

```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```


## Join

Combine tuples on the  
provided predicate




# RA Equivalencies

```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```



# RA Equivalencies

```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```



**Cross Product**  
Same intuition from set theory

# Takeaways

- We can describe relationships between tables with keys and foreign keys
- Different joining techniques can be used to achieve particular goals
- Our SQL toolbox is growing!
  - Not just reading and filtering data anymore
  - Starting to answer complex questions