



1

# Outline

- A Little Extra SQL
- Key Recap
- Introduce Foreign Keys
- Introduce Joins
- Demo in Sqlite

# A Little Extra SQL

- ORDER BY – Orders result tuples by specified attributes (default ascending)

```
SELECT P.UserID, P.Name, P.Salary  
      FROM Payroll AS P  
ORDER BY P.Name ASC
```



Default

```
SELECT P.UserID, P.Name, P.Salary  
      FROM Payroll AS P  
ORDER BY P.Salary DESC
```

# A Little Extra SQL

- **ORDER BY** – Orders result tuples by specified attributes (default ascending)

```
SELECT P.UserID, P.Name, P.Salary  
FROM Payroll AS P  
ORDER BY P.Salary, P.Name;
```

UserID	Name	Salary
123	Jack	50000
345	Allison	50000
567	Magda	90000
789	Dan	100000



UserID	Name	Salary
345	Allison	50000
123	Jack	50000
567	Magda	90000
789	Dan	100000

# A Little Extra SQL

- **DISTINCT** – Deduplicates result tuples
- Data exploration:  
“What are the possible jobs in this dataset?”

```
SELECT DISTINCT Job  
FROM Payroll;
```

Job
TA
Prof

# A Little Extra SQL

- **DISTINCT** – Deduplicates result tuples

```
SELECT P.Job  
      FROM Payroll AS P  
      WHERE P.Salary > 70000;
```

Job
Prof
Prof

# A Little Extra SQL

- **DISTINCT** – Deduplicates result tuples

```
SELECT P.Job  
      FROM Payroll AS P  
      WHERE P.Salary > 70000;
```

Job
Prof
Prof

```
SELECT DISTINCT P.Job  
      FROM Payroll AS P  
      WHERE P.Salary > 70000;
```

Job
Prof

# Preview!

- Data exploration:

“How many people are in this dataset?”



# Preview!

- Data exploration:

“How many people are in this dataset?”

```
SELECT COUNT (*)  
FROM Payroll;
```

COUNT(*)
----------

4
---

# Preview!

- Data exploration:

“How many people are in this dataset?”

```
SELECT COUNT (*) AS num_people  
FROM Payroll;
```

num_people
4

to rename the column

# Recap - Keys

## Key

A **Key** is one or more attributes that **uniquely** identify a row.

Good candidate  
for a key

Definitely not a key

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Recap - Keys

Unique Identifier

```
CREATE TABLE Payroll (  
  UserID INT PRIMARY KEY,  
  Name VARCHAR(100),  
  Job VARCHAR(100),  
  Salary INT);
```

Payroll(UserId, Name, Job, Salary)

# Recap - Keys

```
CREATE TABLE Payroll (  
    UserID INT,  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT,  
    PRIMARY KEY (UserId);
```

Can also define the  
PK on a new line

Payroll(UserId, Name, Job, Salary)

# Recap - Multi-Attribute Keys

Sometimes no single attribute is unique, but combinations of attributes are a unique key for the table.

Must use the PK definition on a new line for multi-attribute keys

```
CREATE TABLE Payroll (  
    Name VARCHAR(100),  
    Job VARCHAR(100),  
    Salary INT,  
    PRIMARY KEY (Name, Job));
```

Here the combination of Name and Job are unique  
e.g. only one "Eden, Professor"  
but some "Eden, TA" or "Ryan, Professor" can exist

Payroll(Name, Job, Salary)

# Foreign Keys

- Databases can hold multiple tables
- How do we capture relationships *between* tables?

**Payroll**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Regist**

UserID	Car
123	Charger
567	Civic
567	Pinto

# Foreign Keys

- Databases can hold multiple tables
- How do we capture relationships *between* tables?

Foreign Key  
UserID


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto



# Foreign Keys

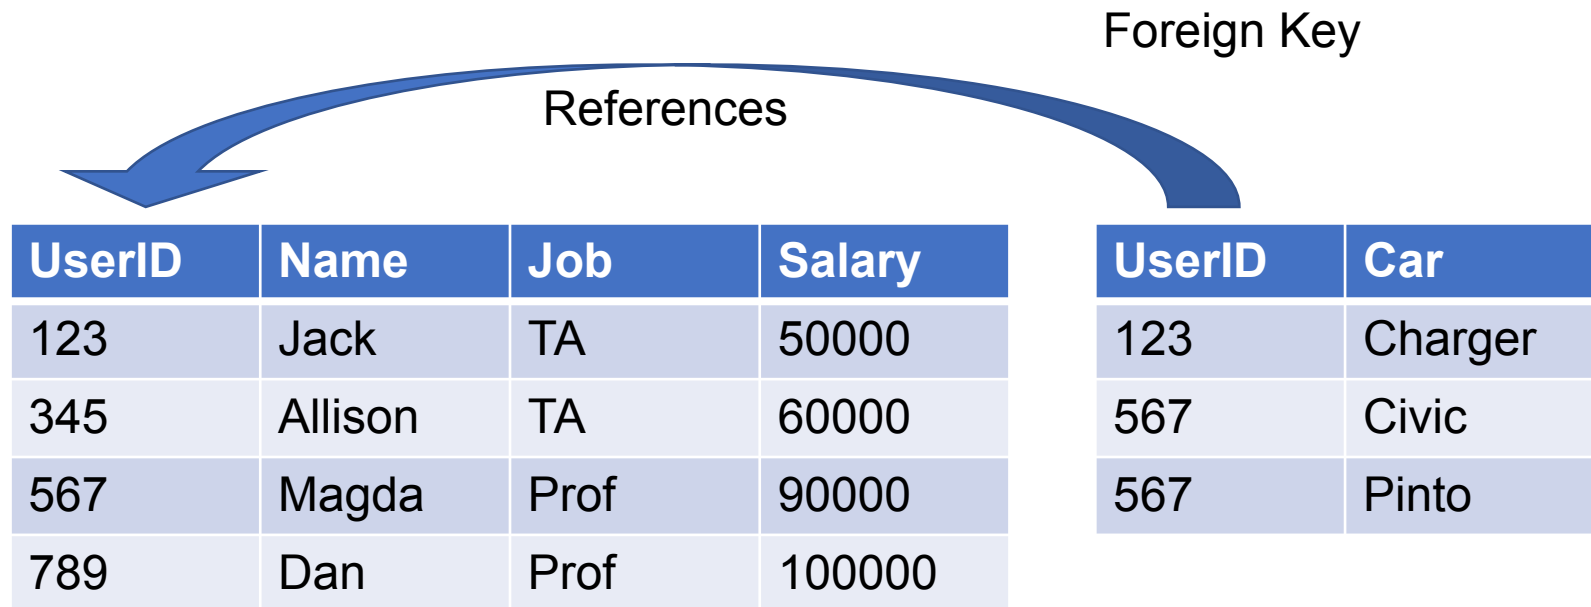
- Databases can hold multiple tables
- How do we capture relationships *between* tables?



# Foreign Keys

## Foreign Key

A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.



# Foreign Keys

## Foreign Key

A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.

Is this valid?

References

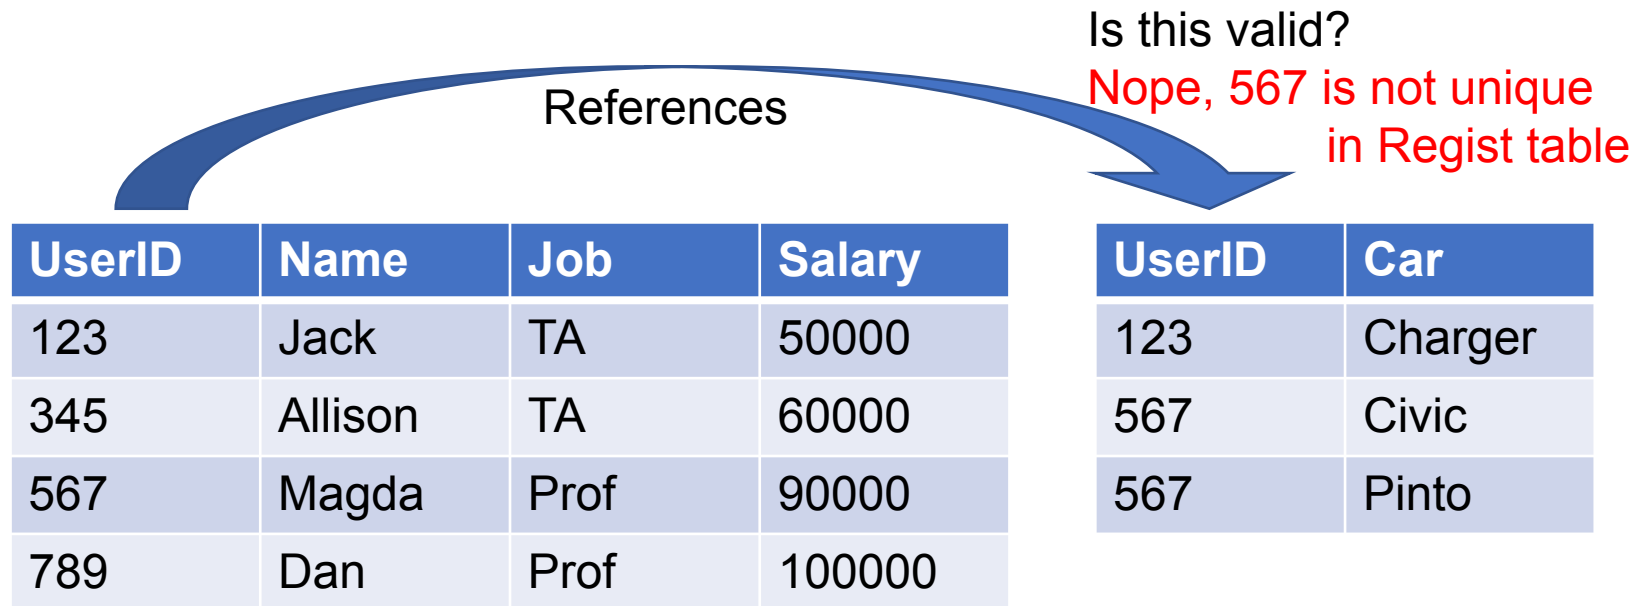
UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

# Foreign Keys

## Foreign Key


A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.



# Foreign Keys

## Foreign Key

A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.



Foreign keys must reference a unique attribute (or set of attributes), almost always the primary key

# Foreign Keys

```
CREATE TABLE Payroll (  
  UserID INT PRIMARY KEY,  
  Name VARCHAR(100),  
  Job VARCHAR(100),  
  Salary INT);
```

Payroll(UserId, Name, Job, Salary)

```
CREATE TABLE Regist (  
  UserID INT,  
  Car VARCHAR(100));
```

Regist(UserId, Car)

# Foreign Keys

```
CREATE TABLE Payroll (  
  UserID INT PRIMARY KEY,  
  Name VARCHAR(100),  
  Job VARCHAR(100),  
  Salary INT);
```

```
CREATE TABLE Regist (  
  UserID INT REFERENCES Payroll,  
  Car VARCHAR(100));
```

Payroll(UserId, Name, Job, Salary)

Regist(UserId, Car)

# Foreign Keys

```
CREATE TABLE Payroll (  
  UserID INT PRIMARY KEY,  
  Name VARCHAR(100),  
  Job VARCHAR(100),  
  Salary INT);
```

```
CREATE TABLE Regist (  
  UserID INT REFERENCES Payroll(UserID),  
  Car VARCHAR(100));
```

or, when attribute name is the same:

```
CREATE TABLE Regist (  
  UserID INT REFERENCES Payroll,  
  Car VARCHAR(100));
```

Payroll(UserId, Name, Job, Salary)

Regist(UserId, Car)



# Foreign Keys

Alternatively, if your foreign key is also more than one attribute:

```
CREATE TABLE Payroll (  
  UserID INT,  
  Name VARCHAR(100),  
  Job VARCHAR(100),  
  Salary INT,  
  PRIMARY KEY(UserID,  
    Name)  
);
```

Payroll(UserID, Name, Job, Salary)

```
CREATE TABLE Regist (  
  UserID INT,  
  Name VARCHAR(100),  
  Car VARCHAR(100),  
  FOREIGN KEY (UserID, Name)  
    REFERENCES Payroll);
```

Regist(UserID, Name, Car)

# Joins

- Foreign keys are able to *describe* a relationship between tables
- Joins are able to realize combinations of data
- Joins do **not** require a foreign key, but often they go together

# Inner Joins

- Bread and butter of SQL queries
  - “Inner join” is often interchangeable with just “join”
- Inner Join syntax:

**Payroll**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Regist**

UserID	Car
123	Charger
567	Civic
567	Pinto

Join Predicate

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R ON P.UserID = R.UserID;
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
ON P.UserID = R.UserID;
```

How do we  
algorithmically  
get our results?

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Register AS R
ON P.UserID = R.UserID;
```

How do we  
algorithmically  
get our results?

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

Compare every possible  
combination and filter  
the results that match

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
ON P.UserID = R.UserID;
```


```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
------	-----


```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```



# Nested-Loop Semantics


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger


```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```



# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto


```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Inner Joins

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Explicit

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
ON P.UserID = R.UserID;
```

Implicit

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

Both of them have the same meaning (for inner joins)

# Inner Joins

```
SELECT P.Name, R.Car  
      FROM Payroll AS P, Regist AS R  
      WHERE P.UserID = R.UserID;
```

- What if we have no join predicate?

```
SELECT P.Name, R.Car  
      FROM Payroll AS P, Regist AS R
```

```
for each row1 in Payroll:  
    for each row2 in Regist:  
        output (row1.Name, row2.Car)
```

- Output every possible pair: “Cross product”



# Outer Joins

Now I want to include everyone, even if they don't drive.

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

# Outer Joins

Now I want to include everyone, even if they don't drive.

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```

# Outer Joins

Now I want to include everyone, even if they don't drive.

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```

# Outer Joins

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Allison	NULL
Magda	Civic
Magda	Pinto
Dan	NULL

# Outer Joins

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Allison	NULL
Magda	Civic
Magda	Pinto
Dan	NULL

NULL is a value placeholder. Depending on context, it may mean unknown, not applicable, etc.

# Outer Joins

- LEFT OUTER JOIN
  - All rows in left table are preserved
- RIGHT OUTER JOIN
  - All rows in right table are preserved
- FULL OUTER JOIN
  - All rows are preserved

# Self-Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      R.Car = 'Civic';
```

# Self-Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      R.Car = 'Civic' AND
      R.Car = 'Pinto' ;
```

Will this work?



# Self-Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      R.Car = 'Civic' AND
      R.Car = 'Pinto';
```

Will this work?

Nope, empty set is  
returned

# Self-Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto
789	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      R.Car = 'Civic' AND
      R.Car = 'Pinto';
```

Discuss with the people around you how you would solve this.

# Self-Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R1.Car
FROM Payroll AS P, Regist AS R1, Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```

# Self-Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

All pairs of cars a person can drive

```
SELECT P.Name, R1.Car
FROM Payroll AS P, Regist AS R1, Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```

# Self-Joins

- When a relation occurs twice in the FROM clause we call it a self-join;
- If we have a self-join, we must use tuple variables (aka table aliases) (why?)

# Self-Joins

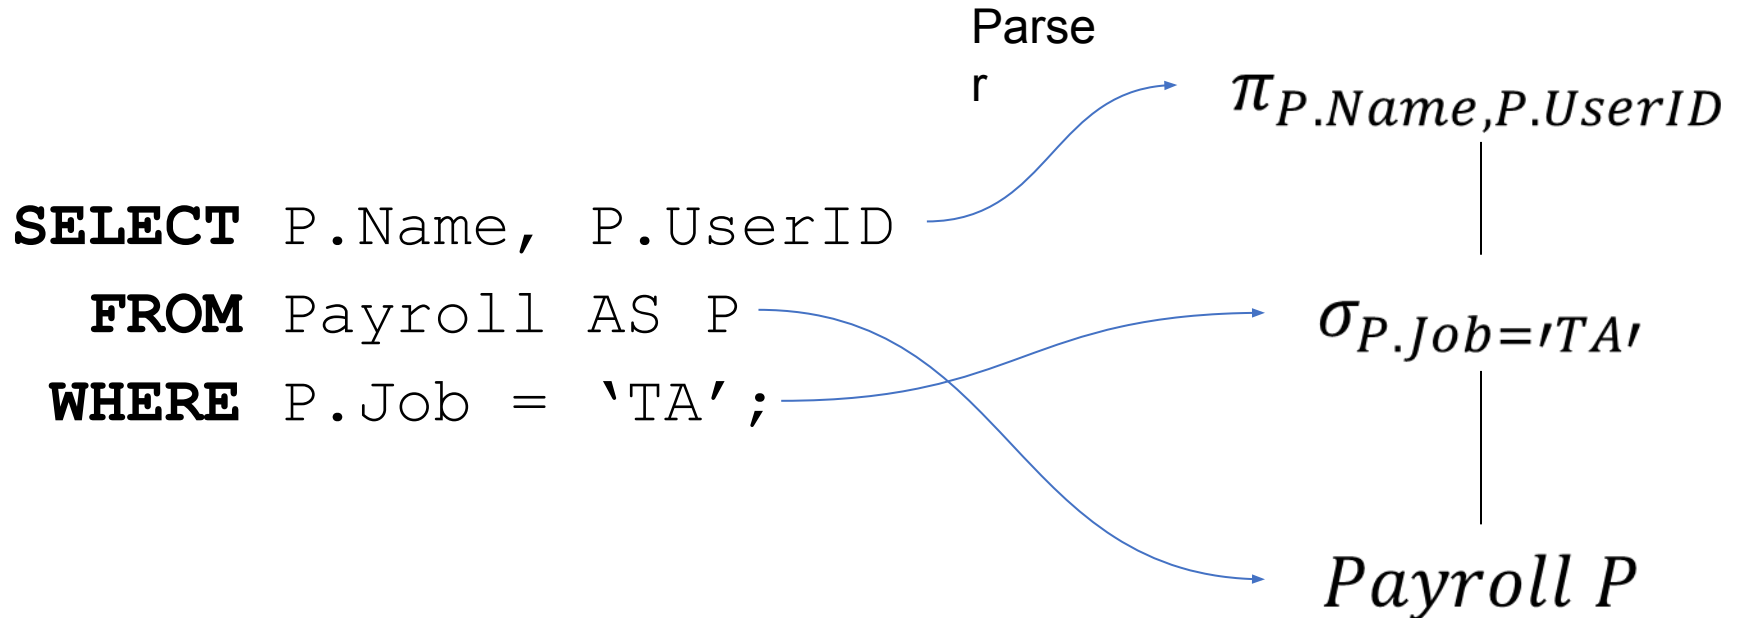
- When a relation occurs twice in the FROM clause we call it a self-join;
- If we have a self-join, we must use tuple variables (aka table aliases)
- Two different tables have an attribute of the same name

# RA Equivalencies

So far we haven't discussed equivalent RA trees.  
But all joins can be parsed directly into a “join tree”

# RA Equivalencies


So far we haven't discussed equivalent RA trees.  
But all joins can be parsed directly into a “join tree”






# RA Equivalencies

So far we haven't discussed equivalent RA trees.  
But all joins can be parsed directly into a “join tree”



# RA Equivalencies

```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```




# RA Equivalencies

```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```


## Join

Combine tuples on the  
provided predicate




# RA Equivalencies

```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```



# RA Equivalencies

```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```



**Cross Product**  
Same intuition from set  
theory

# Takeaways

- We can describe relationships between tables with keys and foreign keys
- Different joining techniques can be used to achieve particular goals
- Our SQL toolbox is growing!
  - Not just reading and filtering data anymore
  - Starting to answer complex questions